
Outer Products for Matrix Multiplication

Technical Report, GSI Technology

Brian Beckman
Technology Fellow
November, 2023

Abstract

Accumulated outer products are much faster than iterated inner products for matrix multiplication.

Iterated Inner Product Versus Accumulated Outer Product

Set Up

Set some array dimensions as mutually co-prime numbers, for generality.

Consider a small matrix of floating-point numbers in the range [0., 1.], inclusive both ends.

```
In[1]:= ClearAll[m, n, k, A, B];  
m = 5; k = 4; n = 7; GCD[m, k, n]  
(A = RandomReal[{0., 1.}, {m, k}]) // MatrixForm
```

```
Out[2]= 1
```

```
Out[3]//MatrixForm=
```

```
( 0.382847  0.0694573  0.627889  0.646139 )  
 ( 0.507911  0.465918  0.380455  0.698418 )  
 ( 0.143523  0.400177  0.545172  0.221861 )  
 ( 0.0517602 0.755199  0.539953  0.536782 )  
 ( 0.533501  0.179299  0.849997  0.79486 )
```

Natively, A is a List of rows with two levels of brackets. Indexing A yields a row vector, with only one level of brackets.

```
In[4]:= A[[2]]
Out[4]= {0.507911, 0.465918, 0.380455, 0.698418}
```

To get a column of A , transpose a row of the transpose. Again, the result has only one level of brackets, so its type as a column vector is implicit. It has exactly the same shape as a row vector. This ambiguity is a defect. However, this defect is commonplace in programming languages.

```
In[5]:= A^T[[2]]^T
Out[5]= {0.0694573, 0.465918, 0.400177, 0.755199, 0.179299}
```

More helper functions

```
In[6]:= ClearAll[row, col];
row[M_, i_] := M[[i]];
col[M_, i_] := M^T[[i]]^T;
```

Make another matrix to play with.

```
In[9]:= (B = RandomReal[{0., 1.}, {k, n}]) // MatrixForm
Out[9]//MatrixForm=
```

$$\begin{pmatrix} 0.827434 & 0.343548 & 0.75827 & 0.433013 & 0.929009 & 0.969251 & 0.259242 \\ 0.0356246 & 0.867495 & 0.336331 & 0.693535 & 0.852175 & 0.872596 & 0.906007 \\ 0.0771809 & 0.926093 & 0.208394 & 0.413351 & 0.0669749 & 0.648475 & 0.115599 \\ 0.955218 & 0.643065 & 0.343872 & 0.182448 & 0.0591317 & 0.672056 & 0.242122 \end{pmatrix}$$

Check Mathematica's built-in matrix product.

```
In[10]:= (A.B) // MatrixForm
Out[10]//MatrixForm=
```

$$\begin{pmatrix} 0.984919 & 1.18877 & 0.666699 & 0.591374 & 0.495118 & 1.27309 & 0.391206 \\ 1.13337 & 1.38014 & 0.861288 & 0.827749 & 0.935678 & 1.61494 & 0.76688 \\ 0.387014 & 1.04401 & 0.433324 & 0.60551 & 0.523987 & 0.990936 & 0.516509 \\ 0.624149 & 1.51815 & 0.590352 & 0.867294 & 0.759551 & 1.42005 & 0.890018 \\ 1.27269 & 1.63715 & 0.915307 & 0.851731 & 0.752351 & 1.75895 & 0.591463 \end{pmatrix}$$

Iterated Inner Product

Iterate inner products -- a school algorithm -- row vectors times column vectors. The following shows row 2 of A *inner-product* col 2 of B to produce element (2, 2) of the product $A.B$.

```
In[11]:= row[A, 2].col[B, 2]
Out[11]=
```

```
1.38014
```

Here is iterated inner product compared for accuracy against the built-in, intrinsic matrix product.

```
In[12]:= On[Assert];
```

```
In[13]:= Assert[Module[{i, j, ab = ConstantArray[0.0, {m, n}]},
  For[i = 1, i ≤ m, i++,
    For[j = 1, j ≤ n, j++,
      ab[[i, j]] = row[A, i].col[B, j]]];
  ab] === A.B]
```

Accumulated Outer Product

Check outer product:

```
In[14]:= Outer[Times, col[A, 2], row[B, 2]] // MatrixForm
Out[14]//MatrixForm=
```

```
( 0.00247439 0.0602539 0.0233606 0.048171 0.0591897 0.0606082 0.0629288 )
( 0.0165981 0.404182 0.156703 0.323131 0.397044 0.406559 0.422125 )
( 0.0142561 0.347152 0.134592 0.277537 0.341021 0.349193 0.362564 )
( 0.0269036 0.655131 0.253997 0.523756 0.643561 0.658983 0.684215 )
( 0.00638743 0.155541 0.0603036 0.12435 0.152794 0.156455 0.162446 )
```

Here is matrix product as accumulated outer product compared against the built-in, intrinsic matrix product. Notice there is only one loop, so this outer-product algorithm should be much faster than the iterated inner-product.

```
In[15]:= Assert[Module[{kk, ab = ConstantArray[0.0, {m, n}]},
  For[kk = 1, kk ≤ k, kk++,
    ab += Outer[Times, col[A, kk], row[B, kk]]];
  ab] === A.B]
```

Helper Functions

Even with these very small matrices, the accumulated outer product is much faster than the iterated inner product.

```

In[16]:= ClearAll[iteratedInnerProduct, accumulatedOuterProduct, builtInProduct];
iteratedInnerProduct[m_, k_, n_, A_, B_] :=
Module[{i, j, ab = ConstantArray[0.0, {m, n}], result, time},
{time, result} = AbsoluteTiming[
For[i = 1, i ≤ m, i++,
For[j = 1, j ≤ n, j++,
ab[[i, j]] = row[A, i].col[B, j]]]; ab];
<|"m" → m, "k" → k, "n" → n, "result" → result,
"time" → Quantity[time, "Seconds"] |>];
accumulatedOuterProduct[m_, k_, n_, A_, B_] :=
Module[{kk, ab = ConstantArray[0.0, {m, n}], result, time},
{time, result} = AbsoluteTiming[
For[kk = 1, kk ≤ k, kk++,
ab += Outer[Times, col[A, kk], row[B, kk]]]; ab];
<|"m" → m, "k" → k, "n" → n, "result" → result,
"time" → Quantity[time, "Seconds"] |>];
builtInProduct[m_, k_, n_, A_, B_] :=
Module[{kk, ab, result, time},
{time, result} = AbsoluteTiming[ab = A.B];
<|"m" → m, "k" → k, "n" → n, "result" → result,
"time" → Quantity[time, "Seconds"] |>];
With[{ab = builtInProduct[m, k, n, A, B],
abi = iteratedInnerProduct[m, k, n, A, B],
aba = accumulatedOuterProduct[m, k, n, A, B]},
Assert[ab["result"] === abi["result"] === aba["result"]];
Print[<|"built-in time" → ab["time"],
"inner time" → abi["time"], "outer time" → aba["time"],
"inner/outer ratio" → N[abi["time"] / aba["time"]],
"outer/built-in ratio" → N[aba["time"] / ab["time"]] |>];

```

```

<|built-in time → 3. × 10-6 s, inner time → 0.000107 s, outer time → 0.000017 s,
inner/outer ratio → 6.29412, outer/built-in ratio → 5.66667 |>

```

Large Matrices

With the following dimensions, iterated inner product takes 20 minutes: too long to wait.

```

In[21]:= GCD[1001, 1114, 993]

```

```

Out[21]=

```

```
1
```

Here are timings for square matrices of every 25 dims from 1 through 376:

In[22]:=

```

ClearAll[timings];
(timings = With[{precision = 1.*^5},
  Module[{timings =
    Table[With[{m = d, k = d, n = d},
      With[{A = RandomReal[{0., 1.}, {m, k}],
        B = RandomReal[{0., 1.}, {k, n}]}],
      <|"dim" → d,
        "built-in" → builtInProduct[m, k, n, A, B],
        "inner" → iteratedInnerProduct[m, k, n, A, B],
        "outer" → accumulatedOuterProduct[m, k, n, A, B] |>
    ]], {d, 1, 400, 25}}],
  Map[Assert[
    Round[#[["built-in"]]["result"], precision] ===
    Round[#[["inner"]]["result"], precision] ===
    Round[#[["outer"]]["result"], precision]
  ] &, timings];
  Map[{#[["dim"]], #[["built-in"]]["time"],
    #[["inner"]]["time"], #[["outer"]]["time"]} &, timings]
]] // MatrixForm

```

Out[23]//MatrixForm=

```

( 1  2. × 10-6 s  0.00001 s  5. × 10-6 s )
26 0.000021 s  0.001819 s  0.000096 s
51 0.000016 s  0.008639 s  0.000257 s
76 0.00026 s   0.034524 s  0.000707 s
101 0.001146 s  0.065588 s  0.00212 s
126 0.00008 s  0.119771 s  0.003335 s
151 0.000101 s  0.178706 s  0.004789 s
176 0.000335 s  0.289882 s  0.006121 s
201 0.000169 s  0.460008 s  0.00777 s
226 0.000206 s  0.647442 s  0.009434 s
251 0.000263 s  0.938203 s  0.011673 s
276 0.00034 s   1.00636 s  0.013534 s
301 0.00044 s   1.40346 s  0.019584 s
326 0.000861 s  1.9191 s   0.021801 s
351 0.000572 s  2.43771 s  0.026874 s
376 0.00086 s   2.7252 s   0.031177 s )

```

A log-linear plot of all the timings:

```
In[24]:= ClearAll[plottableTimings];
plottableTimings[j_] :=
  {col[timings, 1], (Log10@*QuantityMagnitude)[col[timings, j]]}^T
```

```
In[26]:= ListLinePlot[{plottableTimings[2],
  plottableTimings[3], plottableTimings[4]},
  (*ImageSize→Large,*)GridLines→Automatic,
  Frame→True, PlotLegends→{"built-in", "inner", "outer"},
  FrameLabel→
    {"Log10(time[s])", ""}, {"Square Matrix Dimensions", "Running Times"}]]
```

Out[26]=

